

Statement of Research Interest and Bibliography: LLMs and Endangered Language Revitalization

Jared Coleman

Introduction

This document is meant to be an informal, ever-changing collection of interesting papers and resources related to Large Language Models (LLMs) and language revitalization. LLMs have been shown to be remarkably capable at a wide variety of natural language tasks including machine translation, summarizing, question-and-answering, auto-completion, dialog, and more [3]. State-of-the-art LLMs are trained on vast amounts of natural language data from the internet [22] and, as a result, do not perform as well on tasks that involve low/no-resource languages [5, 31]. We refer to languages with very little publicly available bilingual or monolingual corpora as “low-resource” languages and those with *no* publicly available corpora as “no-resource” languages.

Research Questions

In exploring how LLMs might be used for endangered language preservation and revitalization, we have identified the following research questions as some of the most interesting and important:

- How do models “know” language? This is important for understanding how they might be taught new languages from scratch. By taught, I don’t mean fine-tuned or trained (in the ML sense of the word “train”). Rather, I mean *taught* like a human is taught language: through dialog, question and answering, context, and experience.
 - Black-box experimentation: The past few decades have seen many advances in linguistics through creative black-box experiments [2, 9, 25]. Can these be recreated with LLMs? How might the results differ and what might that tell us about how LLMs “know” language?
 - Linguistic Probing: We can perform experiments and “brain-scan” models to see which parts of the underlying network activate to better understand how they work! Interestingly, this has only relatively recently become possible (still with extreme limitations) for humans (via MRI).
 - We care less about whether or not LLMs learn *like* humans and more about understanding how LLMs learn so that we can leverage the knowledge to build useful tools for low/no-resource languages.
- How can we use popular LLM tool-building techniques to create tools for the documentation, preservation, and revitalization of endangered languages?
 - In the context window: few-shot learning, prompt engineering, function calling, etc. We proposed a new approach for low/no-resource language machine translation using a combination of these techniques [7].
 - Tokenization: Can adding tokens for target language words help with natural language tasks?
 - Fine-tuning: with limited data, fine-tuning is difficult.
- How can LLMs be used for foreign language education
 - Ultimately, the goal of endangered language revitalization is to create new *human* speakers.
 - How can LLMs be used effectively in language education? We proposed a new approach for using LLMs as practice partners and tutors for language learning [38].

Useful Tools Enabled by Research

Pursuing the above research questions will guide and enable the development of many practically useful tools for endangered language revitalization. Some of these include:

- Parsing linguistic literature for grammar, vocabulary, etc.
- Summarizing/explaining content for language learners
- Grammar induction
- Auto-completion
- Data sanitization/standardization
- Adaptive data collection: using an LLM to help adjust the questions or queries made to native speakers during data collection to gather the most relevant and useful information.

Special Concerns for Indigenous Communities

When working on language revitalization efforts with indigenous communities, history and context matter. Genocide and forced assimilation [17] have led to the endangerment of many indigenous cultures and languages throughout the United States. At Indian boarding schools, which were established to turn the surviving indigenous population into a servile class, children were forced to abandon their native languages and cultures [16].

Even the more modern and well-intentioned efforts to document and revitalize indigenous languages are not without their own ethical concerns. My tribe, for example, prohibits telling some traditional stories except during the winter. To document these stories and make them publicly available throughout the year would undermine this culturally important tradition. Different indigenous communities have different boundaries and rules for what is appropriate to share and what is not. It is important to respect these boundaries and to work with communities to ensure that the work being done is culturally appropriate and respectful.

Finally, it is imperative that indigenous communities benefit from the work being done to document and revitalize their languages. This means that the tools and resources developed should be made available to the communities in a way that is accessible and useful to them. Another personal example: my great-grandmother was a fluent speaker of our language and so was the subject of a study by the University of California, San Diego Ph.D. student, Evan Norris. His thesis "A Grammar Sketch And Comparative Study Of Eastern Mono" [21], an invaluable resource for our critically endangered language, is locked behind a ProQuest academic paywall and is almost impossible for my family and other tribal members to access.

In our research on using LLMs for endangered language revitalization, we commit to respecting the boundaries and rules of the communities we work with and to making the research output accessible and useful to those communities.

Notes on Papers

This section contains notes, summaries, and thoughts on some of the papers in the bibliography below.

Linguistic Probing

"Does string-based neural MT learn source syntax?" [33] is a very nice introduction to linguistic probing. The following excerpt, in particular, is very helpful in understanding the probing technique in general:

As a simple example, we train an English-French NMT system on 110M tokens of bilingual data (English side). We then take 10K separate English sentences and label their voice as active or passive. We use the learned NMT encoder to convert these sentences into 10k corresponding 1000-dimension encoding vectors. We use 9000 sentences to train a logistic regression model to predict voice using the encoding cell states, and test on the other 1000 sentences. We achieve 92.8% accuracy (Table 2), far above the majority class baseline (82.8%). This means that in reducing the source sentence to a fixed-length vector, the NMT system has decided to store the

voice of English sentences in an easily accessible way. When we carry out the same experiment on an English-English (auto-encoder) system, we find that English voice information is no longer easily accessed from the encoding vector. We can only predict it with 82.7% accuracy, no better than chance. Thus, in learning to reproduce input English sentences, the seq2seq model decides to use the fixed-length encoding vector for other purposes.

So, in general, the idea behind probing is to see whether a model is learning to encode certain known linguistic features as a byproduct of learning to perform a given task. In particular, this example explores how the model is learning to encode the voice of the sentence (active or passive) as a linear combination of the encoding vectors. Note that this approach is *not* capable of telling us about linguistic features that may be encoded in a non-linear way, but it is a good start.

The paper Clark et al. [6] is a very interesting follow-up to the previous paper that applies the probing technique to BERT and explores how different attention heads encode different linguistic features. The paper Manning et al. [18] uses probing techniques to perform experiments that suggest the BERT encoder is learning to encode parse tree distances in its hidden states.

Bibliography

The following bibliography is organized into different categories. Some papers apply to more than one category and therefore appear multiple times.

Our Work

- [7] Jared Coleman et al. *LLM-Assisted Rule Based Machine Translation for Low/No-Resource Languages*. 2024. DOI: [10.48550/arXiv.2405.08997](https://doi.org/10.48550/arXiv.2405.08997).
- [38] Sheng Yu, Jared Coleman, and Bhaskar Krishnamachari. “Chatlang: A Two-Window Approach to Chatbots for Language Learning”. In: (2023). URL: <https://anrg.usc.edu/www/papers/chatlang.pdf>.

Work on LLMs (Large Language Models)

- [3] Sébastien Bubeck et al. “Sparks of Artificial General Intelligence: Early experiments with GPT-4”. In: (2023). DOI: [10.48550/arXiv.2303.08774](https://doi.org/10.48550/arXiv.2303.08774).
- [5] Aakanksha Chowdhery et al. “PaLM: Scaling Language Modeling with Pathways”. In: (2022). DOI: [10.48550/arXiv.2204.02311](https://doi.org/10.48550/arXiv.2204.02311).
- [12] Amr Hendy et al. “How Good Are GPT Models at Machine Translation? A Comprehensive Evaluation”. In: (2023). DOI: [10.48550/arXiv.2302.09210](https://doi.org/10.48550/arXiv.2302.09210).
- [14] Séamus Lankford, Haithem Afli, and Andy Way. “adaptMLLM: Fine-Tuning Multilingual Language Models on Low-Resource Languages with Integrated LLM Playgrounds”. In: *Inf.* 14.12 (2023), p. 638. DOI: [10.3390/INF014120638](https://doi.org/10.3390/INF014120638).
- [15] Patrick S. H. Lewis et al. “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”. In: *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*. Ed. by Hugo Larochelle et al. 2020. URL: <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>.
- [22] OpenAI. “GPT-4 Technical Report”. In: (2023). DOI: [10.48550/arXiv.2303.08774](https://doi.org/10.48550/arXiv.2303.08774).
- [31] Nathaniel R. Robinson et al. “ChatGPT MT: Competitive for High- (but not Low-) Resource Languages”. In: (2023). DOI: [10.48550/arXiv.2309.07423](https://doi.org/10.48550/arXiv.2309.07423).

Work on Low-Resource Languages

- [11] Barry Haddow et al. “Survey of Low-Resource Machine Translation”. In: *Computational Linguistics* 48.3 (Sept. 2022), pp. 673–732. ISSN: 0891-2017. DOI: [10.1162/coli_a_00446](https://doi.org/10.1162/coli_a_00446). eprint: https://direct.mit.edu/coli/article-pdf/48/3/673/2040361/coli_a_00446.pdf.
- [14] Séamus Lankford, Haithem Afli, and Andy Way. “adaptMLLM: Fine-Tuning Multilingual Language Models on Low-Resource Languages with Integrated LLM Playgrounds”. In: *Inf.* 14.12 (2023), p. 638. DOI: [10.3390/INF014120638](https://doi.org/10.3390/INF014120638).
- [27] Surangika Ranathunga et al. “Neural Machine Translation for Low-resource Languages: A Survey”. In: *ACM Comput. Surv.* 55.11 (2023), 229:1–229:37. DOI: [10.1145/3567592](https://doi.org/10.1145/3567592).
- [31] Nathaniel R. Robinson et al. “ChatGPT MT: Competitive for High- (but not Low-) Resource Languages”. In: (2023). DOI: [10.48550/arXiv.2309.07423](https://doi.org/10.48550/arXiv.2309.07423).
- [32] Anton Schäfer et al. *Language Imbalance Can Boost Cross-lingual Generalisation*. 2024. DOI: [10.48550/arXiv.2404.07982](https://doi.org/10.48550/arXiv.2404.07982).
- [37] Daniel Torregrosa et al. “Leveraging Rule-Based Machine Translation Knowledge for Under-Resourced Neural Machine Translation Models”. In: *Proceedings of Machine Translation Summit XVII Volume 2: Translator, Project and User Tracks, MTSummit 2019, Dublin, Ireland, August 19-23, 2019*. Ed. by Mikel L. Forcada et al. European Association for Machine Translation, 2019, pp. 125–133. URL: <https://aclanthology.org/W19-6725/>.

Work on RBMT (Rule-Based Machine Translation)

- [13] Tanmai Khanna et al. “Recent advances in Apertium, a free/open-source rule-based machine translation platform for low-resource languages”. In: *Machine Translation* 35.4 (Dec. 2021), pp. 475–502. ISSN: 1573-0573. DOI: [10.1007/s10590-021-09260-6](https://doi.org/10.1007/s10590-021-09260-6).
- [26] Tommi A Pirinen. “Workflows for kickstarting RBMT in virtually No-Resource Situation”. In: *Proceedings of the 2nd Workshop on Technologies for MT of Low Resource Languages*. Ed. by Alina Karakanta et al. Dublin, Ireland: European Association for Machine Translation, Aug. 2019, pp. 11–16. DOI: [10.18653/V1/W19-6803](https://doi.org/10.18653/V1/W19-6803).
- [37] Daniel Torregrosa et al. “Leveraging Rule-Based Machine Translation Knowledge for Under-Resourced Neural Machine Translation Models”. In: *Proceedings of Machine Translation Summit XVII Volume 2: Translator, Project and User Tracks, MTSummit 2019, Dublin, Ireland, August 19-23, 2019*. Ed. by Mikel L. Forcada et al. European Association for Machine Translation, 2019, pp. 125–133. URL: <https://aclanthology.org/W19-6725/>.

Work on RAG (Retrieval Augmented Generation)

- [15] Patrick S. H. Lewis et al. “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”. In: *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*. Ed. by Hugo Larochelle et al. 2020. URL: <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>.

Work on Embeddings & Semantic Similarity

- [4] Jialun Cao et al. “SemMT: A Semantic-Based Testing Approach for Machine Translation Systems”. In: *ACM Trans. Softw. Eng. Methodol.* 31.2 (Apr. 2022). ISSN: 1049-331X. DOI: [10.1145/3490488](https://doi.org/10.1145/3490488).
- [20] Niklas Muennighoff et al. “MTEB: Massive Text Embedding Benchmark”. In: *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics, EACL 2023, Dubrovnik, Croatia, May 2-6, 2023*. Ed. by Andreas Vlachos and Isabelle Augenstein. Association for Computational Linguistics, 2023, pp. 2006–2029. DOI: [10.18653/V1/2023.EACL-MAIN.148](https://doi.org/10.18653/V1/2023.EACL-MAIN.148).
- [29] Nils Reimers and Iryna Gurevych. “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks”. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Nov. 2019. DOI: [10.48550/arXiv.1908.10084](https://doi.org/10.48550/arXiv.1908.10084).
- [35] Yurun Song, Junchen Zhao, and Lucia Specia. “SentSim: Crosslingual Semantic Evaluation of Machine Translation”. In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2021, Online, June 6-11, 2021*. Ed. by Kristina Toutanova et al. Association for Computational Linguistics, 2021, pp. 3143–3156. DOI: [10.18653/V1/2021.NAACL-MAIN.252](https://doi.org/10.18653/V1/2021.NAACL-MAIN.252).

Work on Linguistic Probing

- [6] Kevin Clark et al. “What Does BERT Look at? An Analysis of BERT’s Attention”. In: *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP, BlackboxNLP@ACL 2019, Florence, Italy, August 1, 2019*. Ed. by Tal Linzen et al. Association for Computational Linguistics, 2019, pp. 276–286. DOI: [10.18653/V1/W19-4828](https://doi.org/10.18653/V1/W19-4828).
- [18] Christopher D. Manning et al. “Emergent linguistic structure in artificial neural networks trained by self-supervision”. In: *Proc. Natl. Acad. Sci. USA* 117.48 (2020), pp. 30046–30054. DOI: [10.1073/PNAS.1907367117](https://doi.org/10.1073/PNAS.1907367117).
- [33] Xing Shi, Inkit Padhi, and Kevin Knight. “Does string-based neural MT learn source syntax?” In: *Proceedings of the 2016 conference on empirical methods in natural language processing*. 2016, pp. 1526–1534. DOI: [10.18653/V1/D16-1159](https://doi.org/10.18653/V1/D16-1159).

Other References

- [1] Jessie Little Doe Baird. “Wopanaak language reclamation program: bringing the language home”. *Journal of Global Indigeneity*, 2(2). 2016. URL: <https://ro.uow.edu.au/jgi/vol2/iss2/7>.
- [2] Mark C Baker. *The atoms of language: The mind’s hidden rules of grammar*. Basic books, 2008.
- [8] Serafin M. Coronel-Molina and Teresa L. McCarty. “Indigenous Language Revitalization in the Americas”. In: 2016. URL: <https://api.semanticscholar.org/CorpusID:217243422>.
- [9] Guy Deutscher. *Through the language glass: Why the world looks different in other languages*. Metropolitan books, 2010.
- [16] K Tsianina Lomawaima and Teresa L McCarty. *"To remain an Indian": Lessons in democracy from a century of Native American education*. Teachers College Press, 2006.
- [17] Benjamin Madley. *An American Genocide: The United States and the California Indian Catastrophe, 1846-1873*. Yale University Press, 2016.
- [19] Christopher Moseley. *Atlas of the World’s Languages in Danger*. Unesco, 2010. ISBN: 978-92-3-104096-2.
- [25] Steven Pinker. *The language instinct: How the mind creates language*. Penguin uK, 2003.
- [34] SIL International. *639 Identifier Documentation: mnr*. Accessed: 11 Mar 2024. 2024. URL: <https://iso639-3.sil.org/code/mnr>.
- [36] Joshua Taylor and Timothy Kochem. “Access and empowerment in digital language learning, maintenance, and revival: a critical literature review”. In: *Diaspora, Indigenous, and Minority Education* 16.4 (2022), pp. 234–245.